


# Chapter 3: Processes

---





# Process Concept

- 👁 An operating system executes a variety of programs:
  - 🍃 Batch system – **jobs**
  - 🍃 Time-shared systems – **user programs** or **tasks**
- 👁 Textbook uses the terms **job** and **process** almost interchangeably
- 👁 **Process** – a program in execution; process execution must progress in sequential fashion
- 👁 Multiple parts
  - 🍃 The program code, also called **text section**
  - 🍃 Current activity including **program counter**, processor registers
  - 🍃 **Stack** containing temporary data
    - ▶ Function parameters, return addresses, local variables
  - 🍃 **Data section** containing global variables
  - 🍃 **Heap** containing memory dynamically allocated during run time





# Process Concept (Cont.)


---

- 👁 Program is **passive** entity stored on disk (**executable file**), process is **active**
  - 🍃 Program becomes process when executable file loaded into memory
- 👁 Execution of program started via GUI mouse clicks, command line entry of its name, etc
- 👁 One program can be several processes
  - 🍃 Consider multiple users executing the same program





# Process in Memory





# Process State


---

- 👁 As a process executes, it changes **state**
  - 🍃 **new**: The process is being created
  - 🍃 **running**: Instructions are being executed
  - 🍃 **waiting**: The process is waiting for some event to occur
  - 🍃 **ready**: The process is waiting to be assigned to a processor
  - 🍃 **terminated**: The process has finished execution





# Diagram of Process State






# Process Control Block (PCB)


Information associated with each process  
(also called **task control block**)

- 👁 Process state – running, waiting, etc
- 👁 Program counter – location of instruction to next execute
- 👁 CPU registers – contents of all process-centric registers
- 👁 CPU scheduling information- priorities, scheduling queue pointers
- 👁 Memory-management information – memory allocated to the process
- 👁 Accounting information – CPU used, clock time elapsed since start, time limits
- 👁 I/O status information – I/O devices allocated to process, list of open files





# CPU Switch From Process to Process







# Threads

---

- 👁 So far, process has a single thread of execution
- 👁 Consider having multiple program counters per process
  - 🍃 Multiple locations can execute at once
    - ▶ Multiple threads of control -> **threads**
- 👁 Must then have storage for thread details, multiple program counters in PCB
- 👁 See next chapter

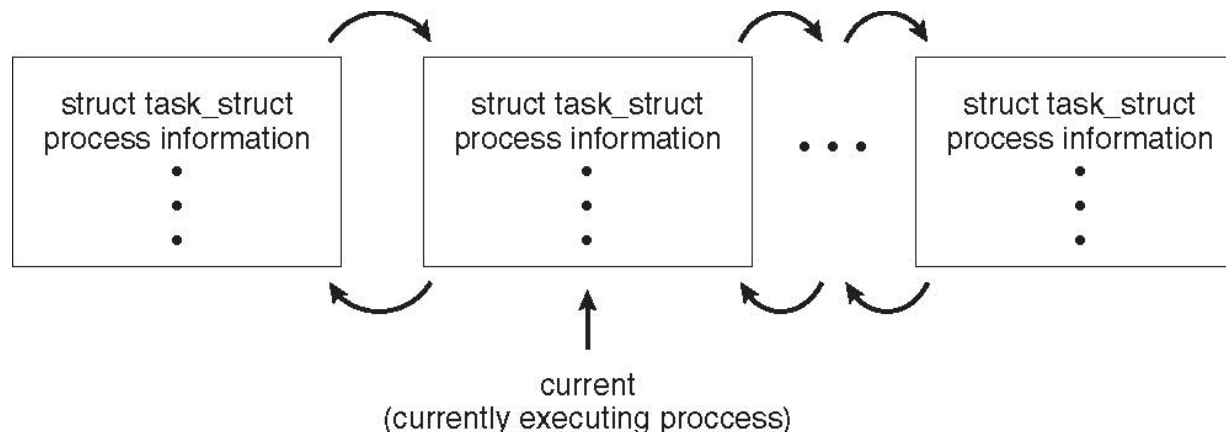




# Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```





# Process Scheduling


---

- 👁️ Maximize CPU use, quickly switch processes onto CPU for time sharing
- 👁️ **Process scheduler** selects among available processes for next execution on CPU
- 👁️ Maintains **scheduling queues** of processes
  - 🍃 **Job queue** – set of all processes in the system
  - 🍃 **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - 🍃 **Device queues** – set of processes waiting for an I/O device
  - 🍃 Processes migrate among the various queues





# Ready Queue And Various I/O Device Queues





# Schedulers

- 👁 **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
  - 🍃 Sometimes the only scheduler in a system
  - 🍃 Short-term scheduler is invoked frequently (milliseconds) ⇒ (must be fast)
- 👁 **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
  - 🍃 Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow)
  - 🍃 The long-term scheduler controls the **degree of multiprogramming**
- 👁 Processes can be described as either:
  - 🍃 **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - 🍃 **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- 👁 Long-term scheduler strives for good ***process mix***





# Context Switch

- 👁 When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- 👁 **Context** of a process represented in the PCB
- 👁 Context-switch time is overhead; the system does no useful work while switching
  - 🍃 The more complex the OS and the PCB => the longer the context switch
- 👁 Time dependent on hardware support
  - 🍃 Some hardware provides multiple sets of registers per CPU  
=> multiple contexts loaded at once





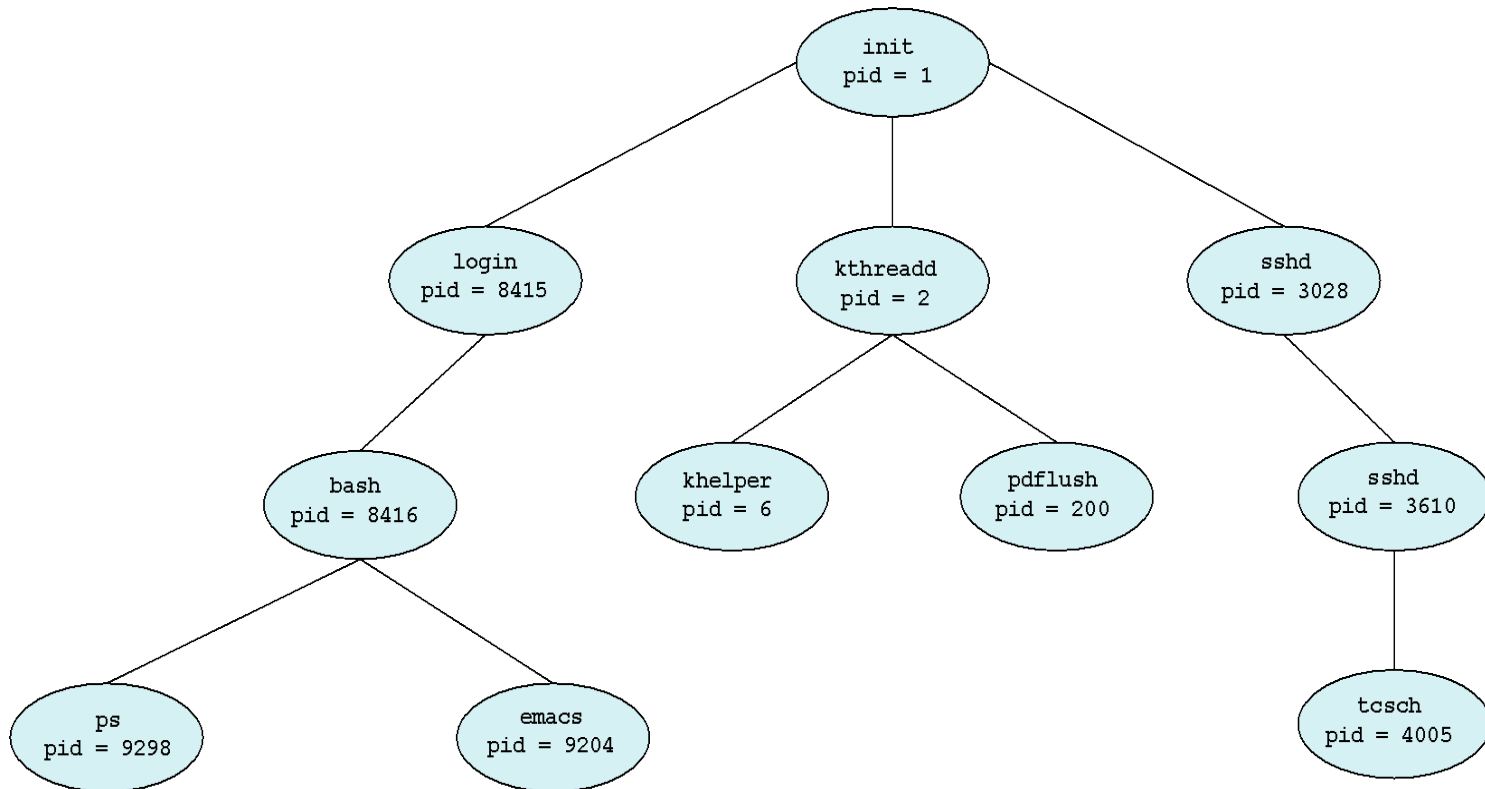
# Process Creation

- 👁 **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- 👁 Generally, process identified and managed via a **process identifier (pid)**
- 👁 Resource sharing options
  - 🍃 Parent and children share all resources
  - 🍃 Children share subset of parent's resources
  - 🍃 Parent and child share no resources
- 👁 Execution options
  - 🍃 Parent and children execute concurrently
  - 🍃 Parent waits until children terminate





# A Tree of Processes in Linux








# Process Creation (Cont.)

- 👁 Address space
  - 🍃 Child duplicate of parent
  - 🍃 Child has a program loaded into it
- 👁 UNIX examples
  - 🍃 **fork()** system call creates new process
  - 🍃 **exec()** system call used after a **fork()** to replace the process' memory space with a new program





# C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





# Process Termination

- 👁 Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
  - 👉 Returns status data from child to parent (via **wait()**)
  - 👉 Process' resources are deallocated by operating system
- 👁 Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
  - 👉 Child has exceeded allocated resources
  - 👉 Task assigned to child is no longer required
  - 👉 The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





# Process Termination

- 👁️ Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
  - 🍃 **cascading termination.** All children, grandchildren, etc. are terminated.
  - 🍃 The termination is initiated by the operating system.
- 👁️ The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process  
**pid = wait(&status);**
- 👁️ If no parent waiting (did not invoke **wait()**) process is a **zombie**
- 👁️ If parent terminated without invoking **wait**, process is an **orphan**





# Multiprocess Architecture – Chrome Browser

- 👁 Many web browsers ran as single process (some still do)
  - 🍃 If one web site causes trouble, entire browser can hang or crash
- 👁 Google Chrome Browser is multiprocess with 3 different types of processes:
  - 🍃 **Browser** process manages user interface, disk and network I/O
  - 🍃 **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
    - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
  - 🍃 **Plug-in** process for each type of plug-in





# Interprocess Communication

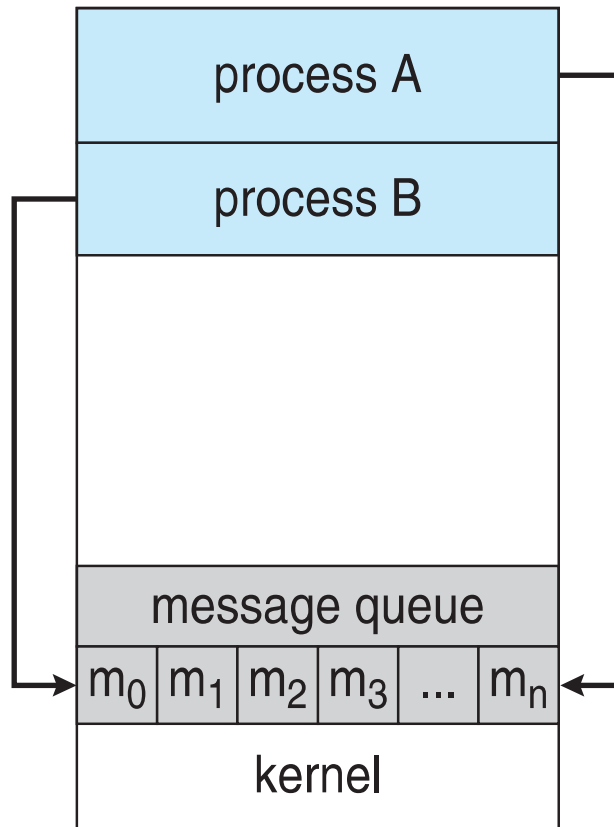
- 👁 Processes within a system may be ***independent*** or ***cooperating***
- 👁 Cooperating process can affect or be affected by other processes, including sharing data
- 👁 Reasons for cooperating processes:
  - 🍃 Information sharing
  - 🍃 Computation speedup
  - 🍃 Modularity
  - 🍃 Convenience
- 👁 Cooperating processes need **interprocess communication (IPC)**
- 👁 Two models of IPC
  - 🍃 **Shared memory**
  - 🍃 **Message passing**






# Communications Models

(a) Message passing. (b) shared memory.



(a)



(b)





# Cooperating Processes

---

- 👁 **Independent** process cannot affect or be affected by the execution of another process
- 👁 **Cooperating** process can affect or be affected by the execution of another process
- 👁 Advantages of process cooperation
  - 🍃 Information sharing
  - 🍃 Computation speed-up
  - 🍃 Modularity
  - 🍃 Convenience







# Producer-Consumer Problem

---

- 👁️ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - 🍃 **unbounded-buffer** places no practical limit on the size of the buffer
  - 🍃 **bounded-buffer** assumes that there is a fixed buffer size





# Bounded-Buffer – Shared-Memory Solution

## 👁 Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

## 👁 Solution is correct, but can only use BUFFER\_SIZE-1 elements





# Bounded-Buffer – Producer

---

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





# Bounded Buffer – Consumer

---

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```





# Interprocess Communication – Shared Memory

---

- 👁 An area of memory shared among the processes that wish to communicate
- 👁 The communication is under the control of the users processes not the operating system.
- 👁 Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- 👁 Synchronization is discussed in great details in Chapter 5.



# End of Chapter 3

---

